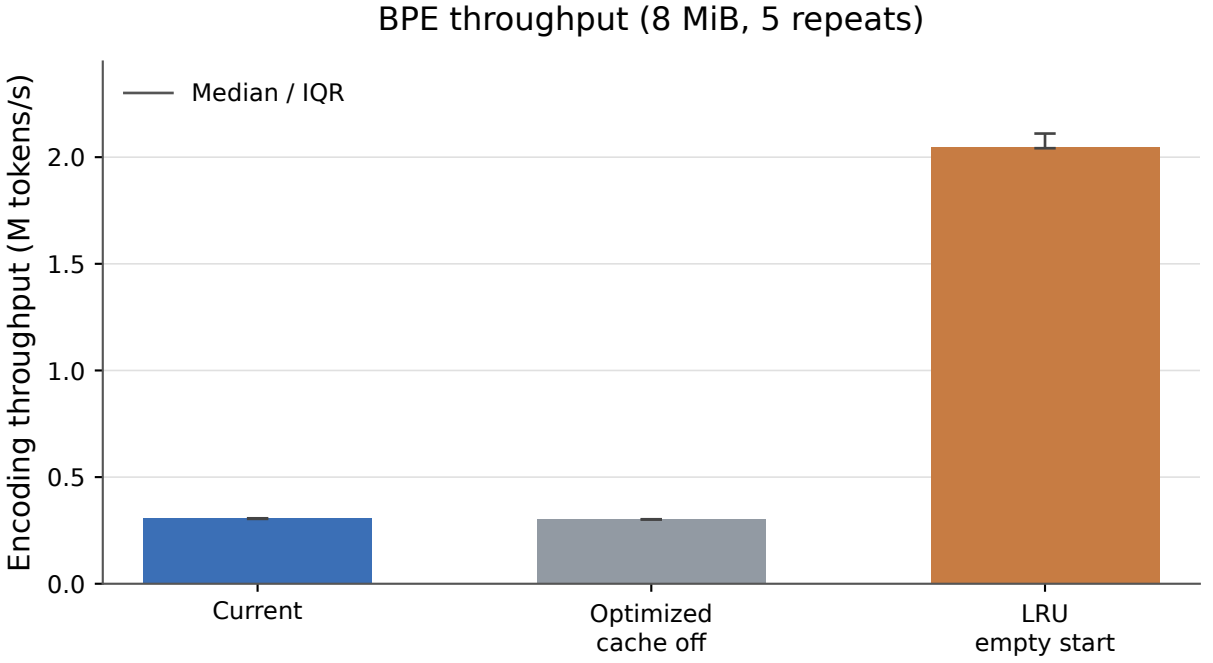
BPE throughput (8 MiB, 5 repeats)
Median / IQR
Encoding throughput (M tokens/s)
0.0
0.5
1.0
1.5
2.0
Current
Optimized
cache off
LRU
empty start

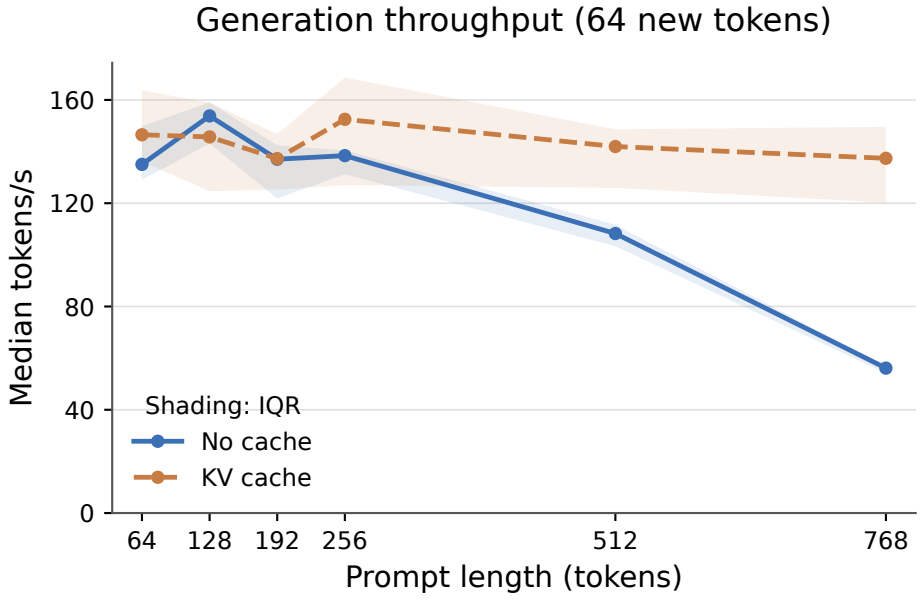
Generation throughput (64 new tokens)
Median tokens/s
160
120
80
40
0
Shading: IQR
No cache
KV cache
64
128
192
256
512
768
Prompt length (tokens)

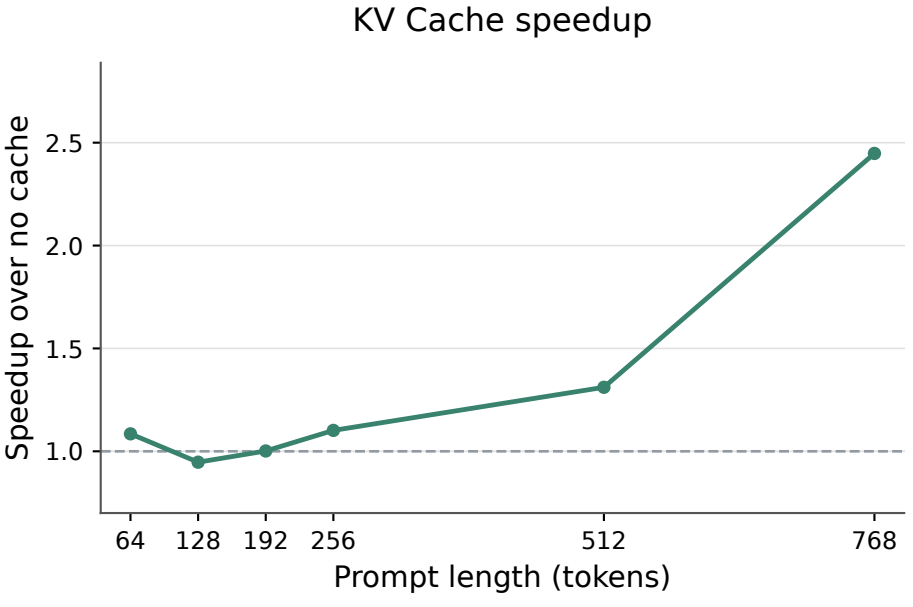
KV Cache speedup
Speedup over no cache
2.5
2.0
1.5
1.0
64
128
192
256
512
768
Prompt length (tokens)

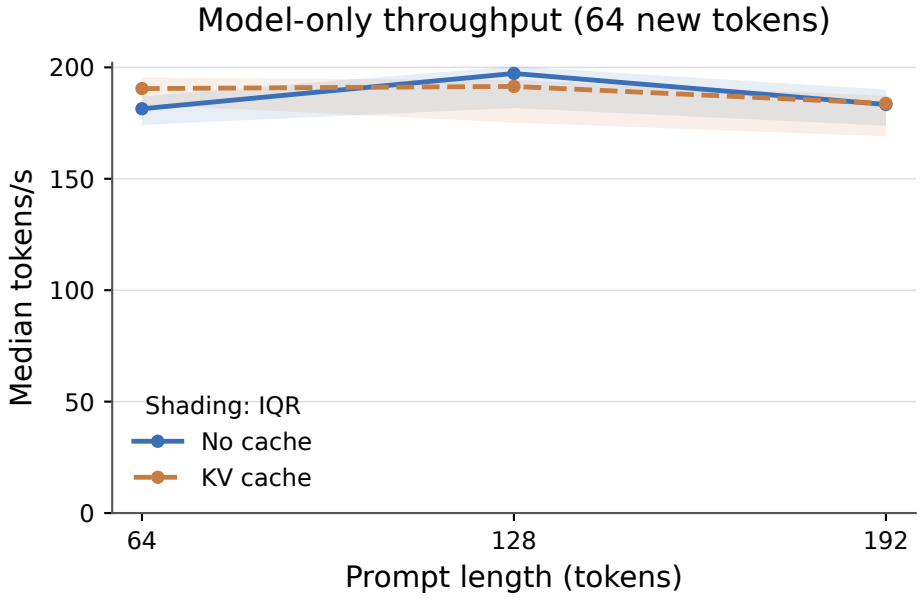
Model-only throughput (64 new tokens)
Median tokens/s
200
150
100
50
0
Shading: IQR
No cache
KV cache
64
128
192
Prompt length (tokens)

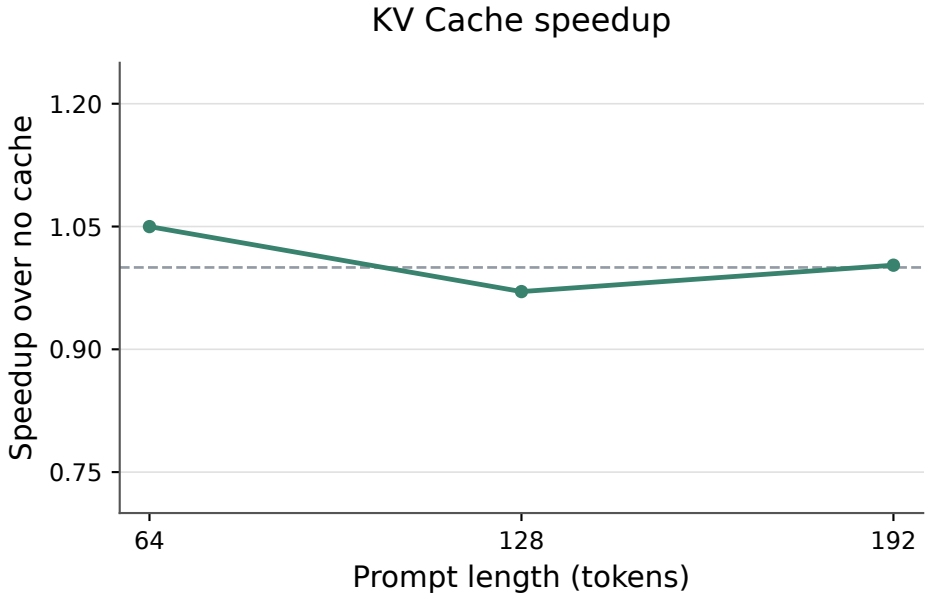
KV Cache speedup
Speedup over no cache
1.20
1.05
0.90
0.75
64
128
192
Prompt length (tokens)